

# CME213 Spring 2023 Homework 2

Ines Dormoy – idormoy@stanford.edu

April 25, 2023

By turning in this assignment, I agree by the Stanford honor code and declare that all of this is my own work.

---

## Problem 1

```
1 Running main() from ./googletest-main/googletest/src/gtest_main.cc
2 [=====] Running 1 test from 1 test suite.
3 [-----] Global test environment set-up.
4 [-----] 1 test from testQ1
5 [ RUN      ] testQ1.test
6 Parallel
7 Sum Even: 757361650
8 Sum Odd: 742539102
9 Time: 0.0120495
10 Serial
11 Sum Even: 757361650
12 Sum Odd: 742539102
13 Time: 0.103744
14 [      OK ] testQ1.test (301 ms)
15 [-----] 1 test from testQ1 (302 ms total)
16
17 [-----] Global test environment tear-down
18 [=====] 1 test from 1 test suite ran. (302 ms total)
19 [  PASSED  ] 1 test.
```

## Problem 2

```
1      Running main() from ./googletest-main/googletest/src/gtest_main.cc
2 [=====] Running 7 tests from 1 test suite.
3 [-----] Global test environment set-up.
4 [-----] 7 tests from testQ2
5 [ RUN      ] testQ2.test1
6 [       OK ] testQ2.test1 (25 ms)
7 [ RUN      ] testQ2.test2
8 [       OK ] testQ2.test2 (2 ms)
9 [ RUN      ] testQ2.test3
10 [      OK ] testQ2.test3 (0 ms)
11 [ RUN      ] testQ2.test4
12 [       OK ] testQ2.test4 (2 ms)
13 [ RUN      ] testQ2.test5
14 [       OK ] testQ2.test5 (6 ms)
15 [ RUN      ] testQ2.serialSortTest
16 Serial Radix Sort: PASS
17 stl: 0.228604
18 serial radix: 0.0880066
19 [       OK ] testQ2.serialSortTest (455 ms)
20 [ RUN      ] testQ2.parallelSortTest
21 Parallel Radix Sort: PASS
22 stl: 0.225698
23 parallel radix: 0.027922
24 [       OK ] testQ2.parallelSortTest (390 ms)
25 [-----] 7 tests from testQ2 (883 ms total)
26
27 [-----] Global test environment tear-down
28 [=====] 7 tests from 1 test suite ran. (883 ms total)
29 [ PASSED ] 7 tests.
```

## Problem 2 Q6

Threads	Blocks	Timing									
	1	2	4	8	12	16	24	32	40	48	
1	0.047	0.046	0.046	0.052	0.052	0.052	0.057	0.062	0.064	0.086	
2	0.045	0.032	0.034	0.030	0.036	0.035	0.038	0.039	0.045	0.057	
3	0.048	0.027	0.021	0.020	0.029	0.024	0.029	0.030	0.034	0.048	
4	0.059	0.026	0.018	0.014	0.018	0.016	0.021	0.022	0.026	0.041	
5	0.055	0.033	0.028	0.019	0.026	0.024	0.030	0.029	0.033	0.049	
6	0.055	0.033	0.024	0.019	0.019	0.019	0.026	0.027	0.031	0.046	
7	0.052	0.035	0.029	0.031	0.025	0.022	0.027	0.030	0.034	0.049	
8	0.053	0.052	0.037	0.025	0.027	0.023	0.035	0.026	0.032	0.049	
9	0.054	0.056	0.041	0.028	0.025	0.021	0.027	0.028	0.031	0.049	
10	0.054	0.057	0.034	0.031	0.026	0.021	0.024	0.028	0.032	0.049	
11											
12											
13											

We observe that with only one block (first column), it is useless to add threads as those threads are not used. Adding those threads is even a waste of time as we spend time creating unused threads. Therefore, the running time increases in the first column when adding threads.

In the second column, we observe that adding a second thread induces a speedup as we use 2 blocks. Again, the runtime increases as we add unused threads.

More globally, the runtime decreases along the diagonal when we add blocks and threads. However, after  $n\_threads = 8$ , we don't observe more speedup (this is because of how the icme-gpu cluster is built). We still observe that the running time increases under this diagonal ( $n\_threads > n\_blocks$ ).

Having  $n\_blocks > n\_threads$  does keep the runtime equivalent when the number of blocks is not too big, but increases runtime if there are too many blocks in comparison to the number of threads (we waste time in the block-threads assignments).

Overall, the best performance is observed with  $n\_threads = n\_blocks$  and  $n\_threads = 8$ .